

Homework 3: Supervised Learning

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```


Exercise 1: Logistic Regression on a Toy 2D Dataset

1. Generate two Gaussian clusters in R^2 and associate them with a class depending on which cluster each point lies on, e.g.:
 - Class 0 centered at $(-2, -2)$ with variance 1,
 - Class 1 centered at $(2, 2)$ with variance 0.5.
2. Plot the dataset in 2D using `plt.scatter` so that each cluster is colored according to its class.

```
n_class_0_samples = 100
n_class_1_samples = 100

class_0 = np.random.normal(0, 1, size=(n_class_0_samples, 2)) +
np.array([-2, -2])
class_1 = np.random.normal(0, 0.5, size=(n_class_1_samples, 2)) +
np.array([2, 2])

plt.scatter(class_0[:, 0], class_0[:, 1], color='blue', label='Class
0')
plt.scatter(class_1[:, 0], class_1[:, 1], color='red', label='Class
1')
plt.title('Synthetic Dataset')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.show()
```



1. Implement logistic regression **from scratch** as did during class:

$$f_{\theta}(x) = \sigma(\theta^T x), \ell(\theta; x, y) = -[y \log f_{\theta}(x) + (1 - y) \log(1 - f_{\theta}(x))].$$

```
def sigmoid(t):
    return 1 / (1 + np.exp(-t))

def l(Theta, X, Y):
    Y_hat = sigmoid(X @ Theta)
    Y_hat = np.clip(Y_hat, 1e-15, 1 - 1e-15) # Avoid log(0)
    return - np.sum(Y * np.log(Y_hat) + (1-Y) * np.log(1 - Y_hat))

def grad_l(Theta, X, Y):
    Y_hat = sigmoid(X @ Theta)
    Y_hat = np.clip(Y_hat, 1e-15, 1 - 1e-15) # Avoid log(0)
    return X.T @ (Y_hat - Y)

def accuracy(Theta, X, Y):
    preds = sigmoid(X @ Theta) >= 0.5
    return (preds == Y).mean()

def SGD(l, grad_l, X, Y, Theta0, lr=1e-2, batch_size=32, epochs=10):
    epoch_loss = []
```

```

epoch_acc = []

Theta = Theta0
for epoch in range(epochs):
    N = len(X)
    shuffle_idx = np.random.permutation(N)

    for start in range(0, N, batch_size):
        batch_idx = shuffle_idx[start:start+batch_size]
        grad = grad_l(Theta, X[batch_idx], Y[batch_idx])
        Theta -= lr * grad

    epoch_loss.append(l(Theta, X, Y))
    epoch_acc.append(accuracy(Theta, X, Y))
    # Update the loss_val list

    return Theta, epoch, epoch_loss, epoch_acc

X = np.vstack((class_0, class_1))
X = np.hstack((np.ones((X.shape[0], 1))), X))
Y = np.hstack((np.zeros(n_class_0_samples),
np.ones(n_class_1_samples)))

print("First 5 samples of X:\n", X[:5])
print("\n")
print("First 5 samples of Y:\n", Y[:5])

print("\n X shape:", X.shape, "Y shape:", Y.shape)

First 5 samples of X:
[[ 1.          -1.03122138 -2.94913065]
 [ 1.          -1.75910219 -2.54883645]
 [ 1.          -1.61269596 -2.9710488 ]
 [ 1.          -2.10609275 -1.53565731]
 [ 1.          -1.99196496 -1.18601649]]

First 5 samples of Y:
[0. 0. 0. 0. 0.]

X shape: (200, 3) Y shape: (200,)

```

1. Train it using simple Gradient Descent on the full dataset. **Note:** the computation of $\nabla L(\Theta; X, Y)$ for this choice of ℓ is given in the teaching note.

```

Theta0 = np.zeros(X.shape[1])
batch_size = X.shape[0]
Theta, epoch, epoch_loss, epoch_acc = SGD(l, grad_l, X, Y,
Theta0=Theta0, lr=0.1, batch_size=batch_size, epochs=10)

```


1. Visualize the learned **decision boundary**:

- plot the line $\{\Theta^T x = 0\}$,
- overlay with the dataset.

```
plt.scatter(class_0[:, 0], class_0[:, 1], color='blue', label='Class 0')
plt.scatter(class_1[:, 0], class_1[:, 1], color='red', label='Class 1')

min_max_of_feature_1 = np.array([min(X[:, 1]) - 1, max(X[:, 1]) + 1])
feature_2 = -(Theta[0] + Theta[1] * min_max_of_feature_1) / Theta[2]
plt.plot(min_max_of_feature_1, feature_2, label='Decision Boundary',
color='green')

plt.title('Synthetic Dataset')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.show()
```



1. Comment on why the decision boundary is linear.

Exercise 2: SGD on Logistic Regression

Use the same synthetic dataset used in the previous exercise, but train logistic regression using **SGD** with the following choices:

1. Try different batch sizes:
 - $N_{\text{batch}}=1$,
 - $N_{\text{batch}}=10$,
 - $N_{\text{batch}}=N$ (full GD).

```
batch_sizes = [1, 10, X.shape[0]]

results = pd.DataFrame([], columns=['Batch Size', 'Theta', 'Final Loss', 'Final Accuracy', 'Loss History', 'Accuracy History'])

for batch_size in batch_sizes:
    Theta0 = np.zeros(X.shape[1])
    Theta, epoch, epoch_loss, epoch_acc = SGD(l, grad_l, X, Y,
    Theta0=Theta0, lr=0.1, batch_size=batch_size, epochs=50)
    results.loc[len(results)] = [batch_size, Theta, epoch_loss[-1],
    epoch_acc[-1], epoch_loss, epoch_acc]
```

```
display(results.sort_values(by='Final Accuracy', ascending=False))
```

	Batch Size	Theta
Final Loss \		
0	1	[-0.8930001306206973, 2.930795967301779, 2.474...
0.129199		
1	10	[-0.8957492503576683, 2.9294798916602773, 2.48...
0.128580		
2	200	[-4.2251669504487595e-05, 21.090594498307944, ...
0.000009		

	Final Accuracy	Loss
History \		
0	1.0	[2.3640703323637613, 1.3875034230363885, 1.027...
1	1.0	[1.855223760635221, 1.2248969128599845, 0.9434...
2	1.0	[8.623084685496744e-06, 8.623073951212868e-06, ...

	Accuracy History
0	[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, ...
1	[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, ...
2	[1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, ...


1. For each setting:
 - Plot the loss vs epoch,

- Plot the classification accuracy vs epoch.

```
n_plots = len(batch_sizes)
n_cols = 2
n_rows = (n_plots + 1) // n_cols
fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 4 *
n_rows))
axes = axes.flatten()
for index, row in results.iterrows():
    batch_size = row['Batch Size']
    loss_history = row['Loss History']
    acc_history = row['Accuracy History']

    ax = axes[index]
    ax.plot(loss_history, label='Loss', color='blue')
    ax.set_xlabel('Epochs')
    ax.set_ylabel('Loss', color='blue')
    ax2 = ax.twinx()
    ax2.plot(acc_history, label='Accuracy', color='red')
    ax2.set_ylabel('Accuracy', color='red')
    ax.set_title(f'Batch Size: {batch_size}')
    ax.legend(loc='upper left')
    ax2.legend(loc='upper right')

# Hide any unused subplots
for j in range(index + 1, len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```



1. Compare the stability and speed of convergence over the choice of different batch sizes.
2. Why the gradients become noisier for small batches? Why larger batches give smoother curves?

Exercise 3: Evaluation Metrics on a Synthetic Dataset

Using the logistic regression model trained above:

1. Compute predicted probabilities $\hat{y}_i = f_{\theta}(x_i)$.
2. Convert them to binary predictions $\hat{y}_i \in \{0, 1\}$ using threshold 0.5.
3. Compute:
 - Confusion matrix (TP, FP, FN, TN),
 - Accuracy,
 - Precision,
 - Recall,
 - F1-score.

```

def classification_report(Y, Y_hat):
    TP = np.sum((Y == 1) & (Y_hat == 1))
    TN = np.sum((Y == 0) & (Y_hat == 0))
    FP = np.sum((Y == 0) & (Y_hat == 1))
    FN = np.sum((Y == 1) & (Y_hat == 0))

    accuracy = (TP + TN) / len(Y)
    precision = TP / (TP + FP) if (TP + FP) > 0 else 0
    recall = TP / (TP + FN) if (TP + FN) > 0 else 0
    f1_score = 2 * (precision * recall) / (precision + recall) if
(precision + recall) > 0 else 0

    confusion_matrix = np.array([[TP, FP],
                                [FN, TN]])

    report = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-Score': f1_score,
        'TP': TP,
        'TN': TN,
        'FP': FP,
        'FN': FN,
        'Confusion Matrix': confusion_matrix
    }
    return report

classification_reports = pd.DataFrame([], columns=['Batch Size',
'Accuracy', 'Precision', 'Recall', 'F1-
Score', 'TP', 'TN', 'FP', 'FN', 'Confusion Matrix'])
for index, row in results.iterrows():
    batch_size = row['Batch Size']
    Theta = row['Theta']
    Y_hat = (sigmoid(X @ Theta) >= 0.5).astype(int)
    report = classification_report(Y, Y_hat)
    classification_reports.loc[len(classification_reports)] =
[batch_size, report['Accuracy'], report['Precision'],
report['Recall'], report['F1-Score'], report['TP'], report['TN'],
report['FP'], report['FN'], report['Confusion Matrix']]

display(classification_reports)


```


	Batch Size	Accuracy	Precision	Recall	F1-Score	TP	TN	FP	FN
0	1	1.0	1.0	1.0	1.0	100	100	0	0
1	10	1.0	1.0	1.0	1.0	100	100	0	0
2	200	1.0	1.0	1.0	1.0	100	100	0	0



```
Confusion Matrix
0 [[100, 0], [0, 100]]
1 [[100, 0], [0, 100]]
2 [[100, 0], [0, 100]]
```

```
from sklearn.metrics import ConfusionMatrixDisplay

for index, row in classification_reports.iterrows():
    batch_size = row['Batch Size']
    cm = row['Confusion Matrix']
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=[0, 1])
    disp.plot(cmap=plt.cm.Blues)
    plt.title(f'Confusion Matrix (Batch Size: {batch_size})')
    plt.show()
```







1. Modify the threshold to:
 - 0.3,
 - 0.7, and repeat.

```
thresholds = [0.3, 0.7]
classification_reports = pd.DataFrame([], columns=['Threshold', 'Batch Size', 'Accuracy', 'Precision', 'Recall', 'F1-Score', 'TP', 'TN', 'FP', 'FN', 'Confusion Matrix'])
for threshold in thresholds:
    for index, row in results.iterrows():
        batch_size = row['Batch Size']
        Theta = row['Theta']
        Y_hat = (sigmoid(X @ Theta) >= threshold).astype(int)
        report = classification_report(Y, Y_hat)
        classification_reports.loc[len(classification_reports)] =
        [threshold, batch_size, report['Accuracy'], report['Precision'],
        report['Recall'], report['F1-Score'], report['TP'], report['TN'],
        report['FP'], report['FN'], report['Confusion Matrix']]

display(classification_reports.sort_values(by=['Accuracy']))
```

	Threshold	Batch Size	Accuracy	Precision	Recall	F1-Score	TP
TN	FP	\					
0	0	0.3	1	1.0	1.0	1.0	100

100	0							
1		0.3	10	1.0	1.0	1.0	1.0	100
100	0							
2		0.3	200	1.0	1.0	1.0	1.0	100
100	0							
3		0.7	1	1.0	1.0	1.0	1.0	100
100	0							
4		0.7	10	1.0	1.0	1.0	1.0	100
100	0							
5		0.7	200	1.0	1.0	1.0	1.0	100
100	0							


	FN	Confusion Matrix
0	0	[[100, 0], [0, 100]]
1	0	[[100, 0], [0, 100]]
2	0	[[100, 0], [0, 100]]
3	0	[[100, 0], [0, 100]]
4	0	[[100, 0], [0, 100]]
5	0	[[100, 0], [0, 100]]

```


from sklearn.metrics import ConfusionMatrixDisplay

for index, row in classification_reports.iterrows():
    batch_size = row['Batch Size']
    cm = row['Confusion Matrix']
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=[0, 1])
    disp.plot(cmap=plt.cm.Blues)
    plt.title(f'Confusion Matrix (Batch Size: {batch_size}),
Threshold: {row["Threshold"]})')
    plt.show()


```




Confusion Matrix (Batch Size: 10), Threshold: 0.3)





Confusion Matrix (Batch Size: 200), Threshold: 0.3)





Confusion Matrix (Batch Size: 10), Threshold: 0.7)





1. Comment on:
 - How lower thresholds increase recall and lower precision,
 - How higher thresholds increase precision and reduce recall,
 - Why classification metrics depend on the application (as discussed in class).

Exercise 4: Logistic Regression on a Real Dataset

Reproduce the pipeline from the book using the **Pima Indians Diabetes Dataset**, available at <https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database>.

This dataset contains medical measurements and a binary label indicating diabetes diagnosis.

1. Preprocess the data as done in class:
 - Download `diabetes.csv`.
 - Extract features X and labels y .
 - **Standardize the features** (mean 0, variance 1).
Explain why normalization is required, connecting to:
 - conditioning,
 - stable optimization,
 - meaningful gradient magnitudes.
 - Add a bias column of 1s.

```

df0 = pd.read_csv('./diabetes.csv')
df0.shape

(768, 9)

df0.head()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```

X = df0.drop(columns=['Outcome']).values
Y = df0['Outcome'].values

X.shape, Y.shape

((768, 8), (768,))

X_mean, X_std = X.mean(axis=0), X.std(axis=0)
X = (X - X_mean) / X_std

X = np.hstack([np.ones((X.shape[0],1)), X])

```

1. Implement logistic regression from scratch (sigmoid + BCE + gradient), and optimize using **SGD** with:
 - a batch size of 32, a learning rate of 10^{-3} , and 200 epochs.
 - for each epoch, track full-dataset BCE loss and full-dataset accuracy.
 - in a plot, visualize the behavior of Loss vs epoch and Accuracy vs epoch. Comment the results.

```

def sigmoid(t):
    return 1 / (1 + np.exp(-t))

def l(Theta, X, Y):
    Y_hat = sigmoid(X @ Theta)

```

```

    return - np.sum(Y * np.log(Y_hat) + (1-Y) * np.log(1 - Y_hat))

def grad_l(Theta, X, Y):
    Y_hat = sigmoid(X @ Theta)
    return X.T @ (Y_hat - Y)

def accuracy(Theta, X, Y):
    preds = sigmoid(X @ Theta) >= 0.5
    return (preds == Y).mean()

def SGD(l, grad_l, X, Y, Theta0, lr=1e-3, batch_size=32, epochs=200):
    epoch_loss = []
    epoch_acc = []

    Theta = Theta0
    for epoch in range(epochs):
        N = len(X)
        shuffle_idx = np.random.permutation(N)

        for start in range(0, N, batch_size):
            batch_idx = shuffle_idx[start:start+batch_size]
            grad = grad_l(Theta, X[batch_idx], Y[batch_idx])
            Theta -= lr * grad

        epoch_loss.append(l(Theta, X, Y))
        epoch_acc.append(accuracy(Theta, X, Y))
        # Update the loss_val list

    return Theta, epoch, epoch_loss, epoch_acc

```

1. Train the same model using **Adam** (using the formulas from class):

$$\Theta_{k+1} = \Theta_k - \eta \frac{\hat{m}_k}{\sqrt{\hat{v}_k + \epsilon}}$$

with a learning rate of 10^{-3} , a batch size of 32, and 200 epochs. Then, plot SGD vs Adam loss and accuracy curves.

```

Theta0 = np.zeros(X.shape[1])
Theta, epoch, epoch_loss, epoch_acc = SGD(l, grad_l, X, Y,
Theta0=Theta0)

```

1. For each method, evaluate:
 - Final accuracy,
 - Confusion matrix,
 - Precision, Recall, F1

```

# plot loss and accuracy curves
plt.figure(figsize=(12, 5))


```

```

plt.subplot(1, 2, 1)
plt.plot(epoch_loss, label='Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epoch_acc, label='Accuracy', color='orange')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Accuracy Curve')
plt.legend()
plt.show()

```



```

def ADAM(l, grad_l, X, Y, Theta0, lr=1e-3, batch_size=32, epochs=200,
        beta1=0.9, beta2=0.999, eps=1e-8):
    epoch_loss = []
    epoch_acc = []

    Theta = Theta0

    m = np.zeros_like(Theta)
    v = np.zeros_like(Theta)
    t = 0

    for epoch in range(epochs):
        N = len(X)
        shuffle_idx = np.random.permutation(N)

        for start in range(0, N, batch_size):

```

```

        batch_idx = shuffle_idx[start:start+batch_size]
        grad = grad_l(Theta, X[batch_idx], Y[batch_idx])

        t += 1
        m = beta1 * m + (1 - beta1) * grad
        v = beta2 * v + (1 - beta2) * (grad * grad)
        m_hat = m / (1 - beta1**t)
        v_hat = v / (1 - beta2**t)
        Theta -= lr * (m_hat / (np.sqrt(v_hat) + eps))

    epoch_loss.append(l(Theta, X, Y))
    epoch_acc.append(accuracy(Theta, X, Y))
    # Update the loss_val list


    return Theta, epoch, epoch_loss, epoch_acc

Theta0 = np.zeros(X.shape[1])
adam_Theta, adam_epoch, adam_epoch_loss, adam_epoch_acc = ADAM(l,
grad_l, X, Y, Theta0=Theta0)

# Plot SDG vs ADAM loss curves
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(epoch_loss, label='SGD Loss', color='blue')
plt.plot(adam_epoch_loss, label='ADAM Loss', color='orange')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve: SGD vs ADAM')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epoch_acc, label='SGD Accuracy', color='blue')
plt.plot(adam_epoch_acc, label='ADAM Accuracy', color='orange')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Accuracy Curve: SGD vs ADAM')
plt.legend()
plt.show()

```



```
Y_hat_SGD = (sigmoid(X @ Theta) >= 0.5).astype(int)
Y_hat_ADAM = (sigmoid(X @ adam_Theta) >= 0.5).astype(int)
SGD_report = classification_report(Y, Y_hat_SGD)
ADAM_report = classification_report(Y, Y_hat_ADAM)

display(pd.DataFrame([SGD_report, ADAM_report], index=['SGD',
'ADAM']))
```


	Accuracy	Precision	Recall	F1-Score	TP	TN	FP	FN	\
SGD	0.782552	0.739336	0.582090	0.651357	156	445	55	112	
ADAM	0.783854	0.740566	0.585821	0.654167	157	445	55	111	


```
Confusion Matrix
SGD  [[156, 55], [112, 445]]
ADAM  [[157, 55], [111, 445]]
```

```
from sklearn.metrics import ConfusionMatrixDisplay
```

```
disp_SGD =
ConfusionMatrixDisplay(confusion_matrix=SGD_report['Confusion
Matrix'], display_labels=[0, 1])
disp_SGD.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix: SGD')
plt.show()
```

```
disp_ADAM =
ConfusionMatrixDisplay(confusion_matrix=ADAM_report['Confusion
Matrix'], display_labels=[0, 1])
disp_ADAM.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix: ADAM')
plt.show()
```





1. Discuss on which method converge faster, which oscillate more, and how this relates to adaptive learning rates discussed in class.

Optional Extension: From Logistic Regression to a Simple Neural Network

This optional section mirrors the final part of the supervised learning chapter.

Implement a neural network with:

- One hidden layer,
- ReLU activation,
- Sigmoid output layer.

You may reuse the implementation shown in the main text. Then:

1. Train the neural network on the **same** Kaggle dataset.
2. Track:
 - BCE loss vs epoch,
 - Accuracy vs epoch.
3. Compare and discuss the performance against logistic regression.

```

import numpy as np

def sigmoid(z):
    return np.where(z >= 0,
                    1/(1+np.exp(-z)),
                    np.exp(z)/(1+np.exp(z)) # prevent overflow when z is
large negative
                    )

def relu(z):
    return np.maximum(0, z)

def relu_grad(z):
    g = np.zeros_like(z)
    g[z > 0] = 1.0
    return g

def bce_loss(y_hat, y, eps=1e-10):
    y_hat = np.clip(y_hat, eps, 1 - eps)
    return -(y*np.log(y_hat) + (1-y)*np.log(1 - y_hat)).mean()

def train_mlp_adam(X, y, hidden=16, lr=1e-2, epochs=200,
batch_size=32):
    N, d = X.shape

    W1 = np.random.normal(0, np.sqrt(2/d), size=(d, hidden))
    b1 = np.zeros((1, hidden))
    W2 = np.random.normal(0, np.sqrt(2/hidden), size=(hidden, 1))
    b2 = np.zeros((1, 1))

    losses, accs = [], []

    for _ in range(epochs):
        # Shuffle data
        idx = np.arange(len(X))
        np.random.shuffle(idx)

        X = X[idx, :]
        y = y[idx]

        for batch in range(0, N, batch_size):
            Xb, yb = X[batch:batch+batch_size],
y[batch:batch+batch_size]

            # forward
            z1 = Xb @ W1 + b1
            a1 = relu(z1)
            z2 = a1 @ W2 + b2
            y_hat = sigmoid(z2)

```

```

        # backward
        delta2 = (y_hat - yb) / len(yb)
        dW2 = a1.T @ delta2
        db2 = delta2.sum(axis=0, keepdims=True)

        delta1 = (delta2 @ W2.T) * relu_grad(z1)
        dW1 = Xb.T @ delta1
        db1 = delta1.sum(axis=0, keepdims=True)

        # update
        W1 -= lr * dW1
        b1 -= lr * db1
        W2 -= lr * dW2
        b2 -= lr * db2

    # track
    z1 = X @ W1 + b1
    a1 = relu(z1)
    z2 = a1 @ W2 + b2
    yhat_full = sigmoid(z2)

    losses.append(bce_loss(yhat_full, y))
    preds = (yhat_full >= 0.5).astype(int)
    accs.append((preds == y).mean())

    return (W1, b1, W2, b2), np.array(losses), np.array(accs)

Y= Y.reshape(-1,1)

X.shape, Y.shape
((768, 9), (768, 1))


theta_nn_sgd, loss_nn_sgd, acc_nn_sgd = train_mlp_adam(
    X, Y,
    hidden=16,
    lr=1e-3,
    epochs=200,
    batch_size=32,
)
print(f"Neural Network final accuracy: {acc_nn_sgd[-1]*100:.2f}%")

Neural Network final accuracy: 74.22%

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(loss_nn_sgd, label='NN Loss', color='blue')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Neural Network Loss Curve')
plt.legend()

```

```
plt.subplot(1, 2, 2)
plt.plot(acc_nn_sgd, label='NN Accuracy', color='orange')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Neural Network Accuracy Curve')
plt.legend()
plt.show()
```



```
# Plot loss and accuracy curves for Adam Vs SGD VS Neural Network
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(epoch_loss, label='SGD Loss', color='blue')
plt.plot(adam_epoch_loss, label='ADAM Loss', color='orange')
plt.plot(loss_nn_sgd, label='Neural Network Loss', color='green')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve: SGD vs ADAM vs Neural Network')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epoch_acc, label='SGD Accuracy', color='blue')
plt.plot(adam_epoch_acc, label='ADAM Accuracy', color='orange')
plt.plot(acc_nn_sgd, label='Neural Network Accuracy', color='green')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Accuracy Curve: SGD vs ADAM vs Neural Network')
plt.legend()

plt.show()
```

